

CHINMMAY V SHASTRY

AI / GenAI Engineer | Bengaluru, India | 91 63605 84277 | chinmay.v.shastri@gmail.com

LinkedIn: linkedin.com/in/chinmayshastry | GitHub: github.com/ChinmayShastry

PROFESSIONAL SUMMARY

AI Engineer specialising in applied GenAI — hybrid RAG pipelines, agentic LLM workflows, and deep learning systems taken from prototype to live deployment. Delivered end-to-end ML pipelines across NLP and computer vision during a project-based AI engineering internship, and independently built and shipped a portfolio of production-style LLM applications using LangChain, the OpenAI API, and vector databases. Prior five years managing equity and derivatives portfolios brings genuine domain fluency and a business-outcome lens to technical decisions. Available to join immediately.

TECHNICAL SKILLS

Languages & Databases: Python, SQL

GenAI & LLM: LangChain, LangGraph, LlamaIndex, OpenAI API (GPT-4o, GPT-4o-mini), Anthropic Claude API, Google Gemini API, Groq, OpenRouter, HuggingFace, Ollama, Prompt Engineering, Embeddings, Agentic Workflows, Multimodal Extraction

Retrieval & Vector Search: Retrieval-Augmented Generation (RAG), FAISS, ChromaDB, BM25, CrossEncoder Reranking, Semantic Chunking, Hybrid Search, RAGAS Evaluation

Machine Learning & Deep Learning: scikit-learn, TensorFlow, Keras, PyTorch, XGBoost, Random Forest, CNN, Transfer Learning, SHAP

Data & Analysis: Pandas, NumPy, SciPy, Matplotlib, Seaborn, Plotly, Feature Engineering, Time Series, Cohort Analysis

Deployment & Tools: Streamlit, Streamlit Cloud, Gradio, HuggingFace Spaces, Git, GitHub, REST APIs, Docker (basic), AWS EC2 / SageMaker (basic)

Visualization & BI: Power BI, Tableau

Domain: Equity markets, derivatives (F&O), portfolio management, quantitative analysis, financial statement analysis

EXPERIENCE

AI Engineer Intern Oct 2025 – June 2026

Rubix AI Solutions Bengaluru, India

- Designed and deployed end-to-end ML pipelines across computer vision and NLP, delivering three production-style systems: Pneumonia Detection CNN, Traffic Sign Recognition, and Indian Sign Language Recognition.
- Owned the full model lifecycle — data preprocessing, feature engineering, model training, hyperparameter tuning, and evaluation — using Python, scikit-learn, and TensorFlow.
- Authored technical documentation covering performance metrics, business impact, and deployment considerations; collaborated cross-functionally on real-world AI solution design.

Portfolio Manager & Financial Advisor 2019 – 2024

Independent / Freelance Bengaluru, India

- Managed equity and derivatives portfolios for 50+ clients, applying fundamental and technical analysis, valuation, and options strategy to optimise risk-adjusted returns.
- Built data-driven screening and monitoring workflows in Python and Excel to track positions, evaluate trade setups, and quantify downside risk across client portfolios.
- Translated complex quantitative findings into clear, outcome-oriented recommendations for non-technical clients — the same stakeholder-communication skill applied to ML results today.

AI / MACHINE LEARNING PROJECTS

StoxAI — AI-Powered Financial Research Assistant | [GitHub](#) · [Streamlit Live](#)

- Architected a financial intelligence platform using **Python and Streamlit**, enabling automated research across Indian equities, commodities, and market indices.
- Integrated **Yahoo Finance, NewsAPI, and GPT-4o-mini** to synthesize technical indicators and news sentiment into comprehensive, long-form research reporting.
- Engineered a **tool-calling chatbot** for multi-asset comparison and natural language querying of real-time and cached market data.
- Optimized performance via disk-based caching and implemented **safety guardrails** to ensure output compliance and prevent unsolicited investment advice.

DocChat AI — Hybrid RAG Document Intelligence | [GitHub](#) · [Streamlit Live](#)

- Engineered a multimodal document chatbot enabling conversational Q&A over uploaded **PDF, Word, Excel, and CSV** assets.

- Implemented a **hybrid RAG architecture** combining **BM25** and **semantic search** for high-fidelity retrieval and grounded response generation with explicit source citations.
- Developed using **Python, Streamlit, OpenAI, and ChromaDB**; integrated advanced document parsing, LRU caching, and automated hallucination verification for production-style reliability.

Enterprise RAG System — Production Deployment | [GitHub](#)

- Built a modular, production-ready RAG pipeline with RBAC-enforced document access, rotating query logging, LRU caching, and Tenacity-based retry logic for API resilience.
- Added a feedback dashboard for continuous quality monitoring; RBAC design mirrors real data-governance requirements in regulated environments.

LLM-Powered Resume Analyzer — Agentic AI Tool | [GitHub](#) · [Streamlit Live](#)

- Architected an agentic evaluation system with autonomous multi-step reasoning (parse → extract → analyse → output) producing structured, role-specific skill-gap analysis and candidate recommendations.
- Deployed an interactive Streamlit dashboard for real-time scoring and feedback visualisation.

Computer Vision & Deep Learning Portfolio | [GitHub](#)

- Benchmarked a from-scratch CNN against VGG16 and ResNet50 transfer learning for chest X-ray pneumonia detection; ResNet50 reached 88% accuracy and 0.95 AUC, scored on sensitivity and specificity separately rather than accuracy alone.
- Architected an Indian Sign Language Recognition system classifying 24 gestures at 99.87% accuracy on a 4,972-image custom dataset; engineered a 43-class Traffic Sign Classifier using MobileNetV2 and ResNet50 with stratified pipelines for robust generalisation.

Predictive Modeling & Quantitative Analysis

- Designed a Telecom Churn Prediction model using XGBoost and SHAP across a 243K-customer dataset, quantifying cost-benefit trade-offs to drive retention strategy.
- Built a custom OLS-estimated SARIMA model with ADF/ACF/PACF diagnostics for recursive forecasting — framework transfers directly to financial time series and volatility modelling.
- Implemented K-Nearest Neighbours and linear regression from first principles in NumPy — gradient descent, loss tracking and feature scaling written by hand, then validated against scikit-learn.

EDUCATION & PROFESSIONAL TRAINING

AI Engineer Professional Program (Offline) — DataMites 2025

Bachelor of Commerce (B.Com), 81.43% — Gulbarga University Feb 2025

CERTIFICATIONS

AI Engineer Certification — NASSCOM 2026

AI Engineer Certification — IABAC 2026